

Oracle 11g - SQL

Data Manipulating

Objectives

After completing this lesson, you should be able to do the following:

- Describe each data manipulation language (DML) statement
- Insert rows into a table
- Update rows in a table
- Delete rows from a table
- Control transactions
- Read Consistency

Data Manipulation Language

- A DML statement is executed when you:
 - o Add new rows to a table
 - o Modify existing rows in a table
 - o Remove existing rows from a table
- A *transaction* consists of a collection of DML statements that form a logical unit of work.

Adding a New Row to a Table

DEPARTMENTS

| | DEPARTMENT_ID | DEPARTMENT_NAME | MANAGER_ID | LOCATION_ID |
|---|---------------|-----------------|------------|-------------|
| 1 | 10 | Administration | 200 | 1700 |
| 2 | 20 | Marketing | 201 | 1800 |
| 3 | 50 | Shipping | 124 | 1500 |
| 4 | 60 | IT | 103 | 1400 |
| 5 | 80 | Sales | 149 | 2500 |
| 6 | 90 | Executive | 100 | 1700 |
| 7 | 110 | Accounting | 205 | 1700 |
| 8 | 190 | Contracting | (null) | 1700 |

| | | |
|---------------------|-----|------|
| 70 Public Relations | 100 | 1700 |
|---------------------|-----|------|

New
row

Insert new row
into the
DEPARTMENTS table

| | DEPARTMENT_ID | DEPARTMENT_NAME | MANAGER_ID | LOCATION_ID |
|---|---------------|-----------------|------------|-------------|
| 1 | 10 | Administration | 200 | 1700 |
| 2 | 20 | Marketing | 201 | 1800 |
| 3 | 50 | Shipping | 124 | 1500 |
| 4 | 60 | IT | 103 | 1400 |
| 5 | 80 | Sales | 149 | 2500 |
| 6 | 90 | Executive | 100 | 1700 |
| 7 | 110 | Accounting | 205 | 1700 |
| 8 | 190 | Contracting | (null) | 1700 |
| 9 | 190 | Contracting | (null) | 1700 |

INSERT Statement Syntax

- Add new rows to a table by using the `INSERT` statement:

```
INSERT INTO  table [(column [, column...])]  
VALUES      (value [, value...]);
```

- With this syntax, only one row is inserted at a time.

Inserting New Rows

- Insert a new row containing values for each column.
- List values in the default order of the columns in the table.
- Optionally, list the columns in the `INSERT` clause.

```
INSERT INTO departments (department_id,  
                        department_name, manager_id, location_id)  
VALUES (70, 'Public Relations', 100, 1700);  
1 rows inserted
```

- Enclose character and date values in single quotation marks.

Inserting Rows with Null Values

- Implicit method: Omit the column from the column list.

```
INSERT INTO departments (department_id,  
                          department_name )  
VALUES (30, 'Purchasing');  
1 rows inserted
```

- Explicit method: Specify the NULL keyword in the VALUES clause.

```
INSERT INTO departments  
VALUES (100, 'Finance', , );  
1 rows inserted
```

Inserting Special Values

The SYSDATE function records the current date and time.

```
INSERT INTO employees (employee_id,  
                        first_name, last_name,  
                        email, phone_number,  
                        hire_date, job_id, salary,  
                        commission_pct, manager_id,  
                        department_id)  
VALUES  
    (113,  
     'Louis', 'Popp',  
     'LPOPP', '515.124.4567',  
     SYSDATE, 'AC_ACCOUNT', 6900,  
     NULL, 205, 100);
```

1 rows inserted

Inserting Specific Date Values

- Add a new employee.

```
INSERT INTO employees
VALUES      (114,
             'Den', 'Raphealy',
             'DRAPHEAL', '515.127.4561',
             TO_DATE('FEB 3, 1999', 'MON DD, YYYY'),
             'AC_ACCOUNT', 11000, NULL, 100, 30);
1 rows inserted
```

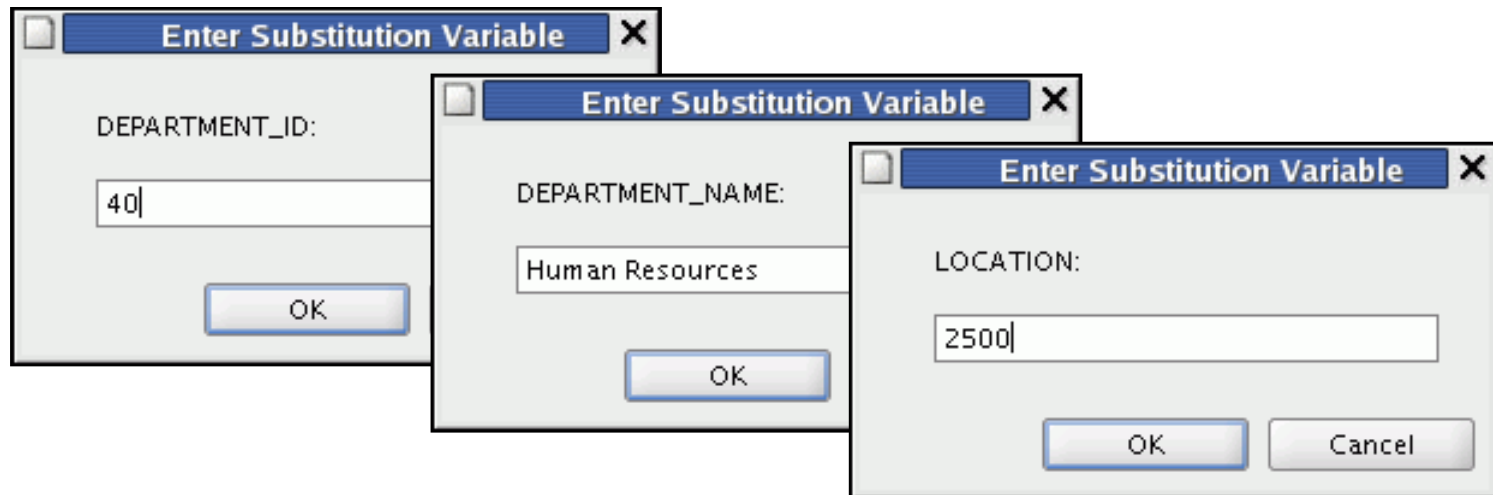
- Verify your addition.

| | EMPLOYEE_ID | FIRST_NAME | LAST_NAME | EMAIL | PHONE_NUMBER | HIRE_DATE | JOB_ID | SALARY | COMMISSION_PCT |
|---|-------------|------------|-----------|----------|--------------|-----------|------------|--------|----------------|
| 1 | 114 | Den | Raphealy | DRAPHEAL | 515.127.4561 | 03-FEB-99 | AC_ACCOUNT | 11000 | (null) |

Creating a Script

- Use & substitution in a SQL statement to prompt for values.
- & is a placeholder for the variable value.

```
INSERT INTO departments
      (department_id, department_name, location_id)
VALUES (&department_id, '&department_name', &location);
```



Copying Rows from Another Table

- Write your `INSERT` statement with a subquery:

```
INSERT INTO sales_reps(id, name, salary, commission_pct)
SELECT employee_id, last_name, salary, commission_pct
FROM   employees
WHERE  job_id LIKE '%REP%';
```

4 rows inserted

- Do not use the `VALUES` clause.
- Match the number of columns in the `INSERT` clause to those in the subquery.

Changing Data in a Table

EMPLOYEES

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | EMAIL | HIRE_DATE | JOB_ID | SALARY | DEPARTMENT_ID | COMMISSION_PCT |
|-------------|------------|-----------|----------|-----------|---------|--------|---------------|----------------|
| 100 | Steven | King | SKING | 17-JUN-87 | AD_PRES | 24000 | 90 | (null) |
| 101 | Neena | Kochhar | NKOCHHAR | 21-SEP-89 | AD_VP | 17000 | 90 | (null) |
| 102 | Lex | De Haan | LDEHAAN | 13-JAN-93 | AD_VP | 17000 | 90 | (null) |
| 103 | Alexander | Hunold | AHUNOLD | 03-JAN-90 | IT_PROG | 9000 | 60 | (null) |
| 104 | Bruce | Ernst | BERNST | 21-MAY-91 | IT_PROG | 6000 | 60 | (null) |
| 107 | Diana | Lorentz | DLORENTZ | 07-FEB-99 | IT_PROG | 4200 | 60 | (null) |
| 124 | Kevin | Mourgos | KMOURGOS | 16-NOV-99 | ST_MAN | 5800 | 50 | (null) |

Update rows in the EMPLOYEES table:

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | EMAIL | HIRE_DATE | JOB_ID | SALARY | DEPARTMENT_ID | COMMISSION_PCT |
|-------------|------------|-----------|----------|-----------|---------|--------|---------------|----------------|
| 100 | Steven | King | SKING | 17-JUN-87 | AD_PRES | 24000 | 90 | (null) |
| 101 | Neena | Kochhar | NKOCHHAR | 21-SEP-89 | AD_VP | 17000 | 90 | (null) |
| 102 | Lex | De Haan | LDEHAAN | 13-JAN-93 | AD_VP | 17000 | 90 | (null) |
| 103 | Alexander | Hunold | AHUNOLD | 03-JAN-90 | IT_PROG | 9000 | 30 | (null) |
| 104 | Bruce | Ernst | BERNST | 21-MAY-91 | IT_PROG | 6000 | 30 | (null) |
| 107 | Diana | Lorentz | DLORENTZ | 07-FEB-99 | IT_PROG | 4200 | 30 | (null) |
| 124 | Kevin | Mourgos | KMOURGOS | 16-NOV-99 | ST_MAN | 5800 | 50 | (null) |

UPDATE Statement Syntax

- Modify existing rows with the `UPDATE` statement:

```
UPDATE      table  
SET         column = value [, column = value, ...]  
[WHERE      condition];
```

- Update more than one row at a time (if required).

Updating Rows in a Table

- Specific row or rows are modified if you specify the `WHERE` clause:

```
UPDATE employees
SET    department_id = 70
WHERE  employee_id = 113;
1 rows updated
```

- All rows in the table are modified if you omit the `WHERE` clause:

```
UPDATE    copy_emp
SET       department_id = 110;
22 rows updated
```

Updating Two Columns with a Subquery

Update employee 114's job and salary to match that of employee 205.

```
UPDATE    employees
SET       job_id    = (SELECT  job_id
                        FROM      employees
                        WHERE     employee_id = 205),
          salary    = (SELECT  salary
                        FROM      employees
                        WHERE     employee_id = 205)
WHERE     employee_id    = 114;
1 rows updated
```

Updating Rows Based on Another Table

Use subqueries in `UPDATE` statements to update rows in a table based on values from another table:

```
UPDATE copy_emp
SET    department_id = (SELECT department_id
                        FROM employees
                        WHERE employee_id = 100)
WHERE  job_id = (SELECT job_id
                 FROM employees
                 WHERE employee_id = 200);

1 rows updated
```


Removing a Row from a Table

DEPARTMENTS

| R# | DEPARTMENT_ID | DEPARTMENT_NAME | R# | MANAGER_ID | R# | LOCATION_ID |
|----|---------------|-----------------|----|------------|----|-------------|
| 1 | 30 | Purchasing | | (null) | | (null) |
| 2 | 40 | Human Resources | | (null) | | 2500 |
| 3 | 10 | Administration | | 200 | | 1700 |
| 4 | 20 | Marketing | | 201 | | 1800 |
| 5 | 50 | Shipping | | 124 | | 1500 |
| 6 | 60 | IT | | 103 | | 1400 |

Delete a row from the DEPARTMENTS table:

| R# | DEPARTMENT_ID | DEPARTMENT_NAME | R# | MANAGER_ID | R# | LOCATION_ID |
|----|---------------|-----------------|----|------------|----|-------------|
| 1 | 40 | Human Resources | | (null) | | 2500 |
| 2 | 10 | Administration | | 200 | | 1700 |
| 3 | 20 | Marketing | | 201 | | 1800 |
| 4 | 50 | Shipping | | 124 | | 1500 |
| 5 | 60 | IT | | 103 | | 1400 |

DELETE Statement

You can remove existing rows from a table by using the `DELETE` statement:

```
DELETE [FROM]   table  
[WHERE         condition];
```

Deleting Rows from a Table

- Specific rows are deleted if you specify the `WHERE`

```
DELETE FROM departments  
WHERE department_name = 'Finance';  
1 rows deleted
```

- All rows in the table are deleted if you omit the

```
DELETE FROM copy_emp;  
22 rows deleted
```

Deleting Rows Based on Another Table

Use subqueries in `DELETE` statements to remove rows from a table based on values from another table:

```
DELETE FROM employees
WHERE department_id =
    (SELECT department_id
     FROM departments
     WHERE department_name
           LIKE '%Public%');

1 rows deleted
```

TRUNCATE Statement

- Removes all rows from a table, leaving the table empty and the table structure intact
- Is a data definition language (DDL) statement rather than a DML statement; cannot easily be undone

```
TRUNCATE TABLE table_name;
```

```
TRUNCATE TABLE copy_emp;
```

Using a Subquery in an INSERT Statement

```
INSERT INTO
    (SELECT employee_id, last_name,
            email, hire_date, job_id, salary,
            department_id
     FROM   employees
     WHERE  department_id = 50)
VALUES (99999, 'Taylor', 'DTAYLOR',
        TO_DATE('07-JUN-99', 'DD-MON-RR'),
        'ST_CLERK', 5000, 50);
```

1 rows inserted

Using a Subquery in an INSERT Statement

Verify the results:

```
SELECT employee_id, last_name, email, hire_date,  
       job_id, salary, department_id  
FROM   employees  
WHERE  department_id = 50;
```

| | EMPLOYEE_ID | LAST_NAME | EMAIL | HIRE_DATE | JOB_ID | SALARY | DEPARTMENT_ID |
|---|-------------|-----------|----------|-----------|----------|--------|---------------|
| 1 | 99999 | Taylor | DTAYLOR | 07-JUN-99 | ST_CLERK | 5000 | 50 |
| 2 | 124 | Mourgos | KMOURGOS | 16-NOV-99 | ST_MAN | 5800 | 50 |
| 3 | 141 | Rajs | TRAJS | 17-OCT-95 | ST_CLERK | 3500 | 50 |
| 4 | 142 | Davies | CDAVIES | 29-JAN-97 | ST_CLERK | 3100 | 50 |
| 5 | 143 | Matos | RMATOS | 15-MAR-98 | ST_CLERK | 2600 | 50 |
| 6 | 144 | Vargas | PVARGAS | 09-JUL-98 | ST_CLERK | 2500 | 50 |

Database Transactions

A database transaction consists of one of the following:

- DML statements that constitute one consistent change to the data
- One DDL statement
- One data control language (DCL) statement

Database Transactions

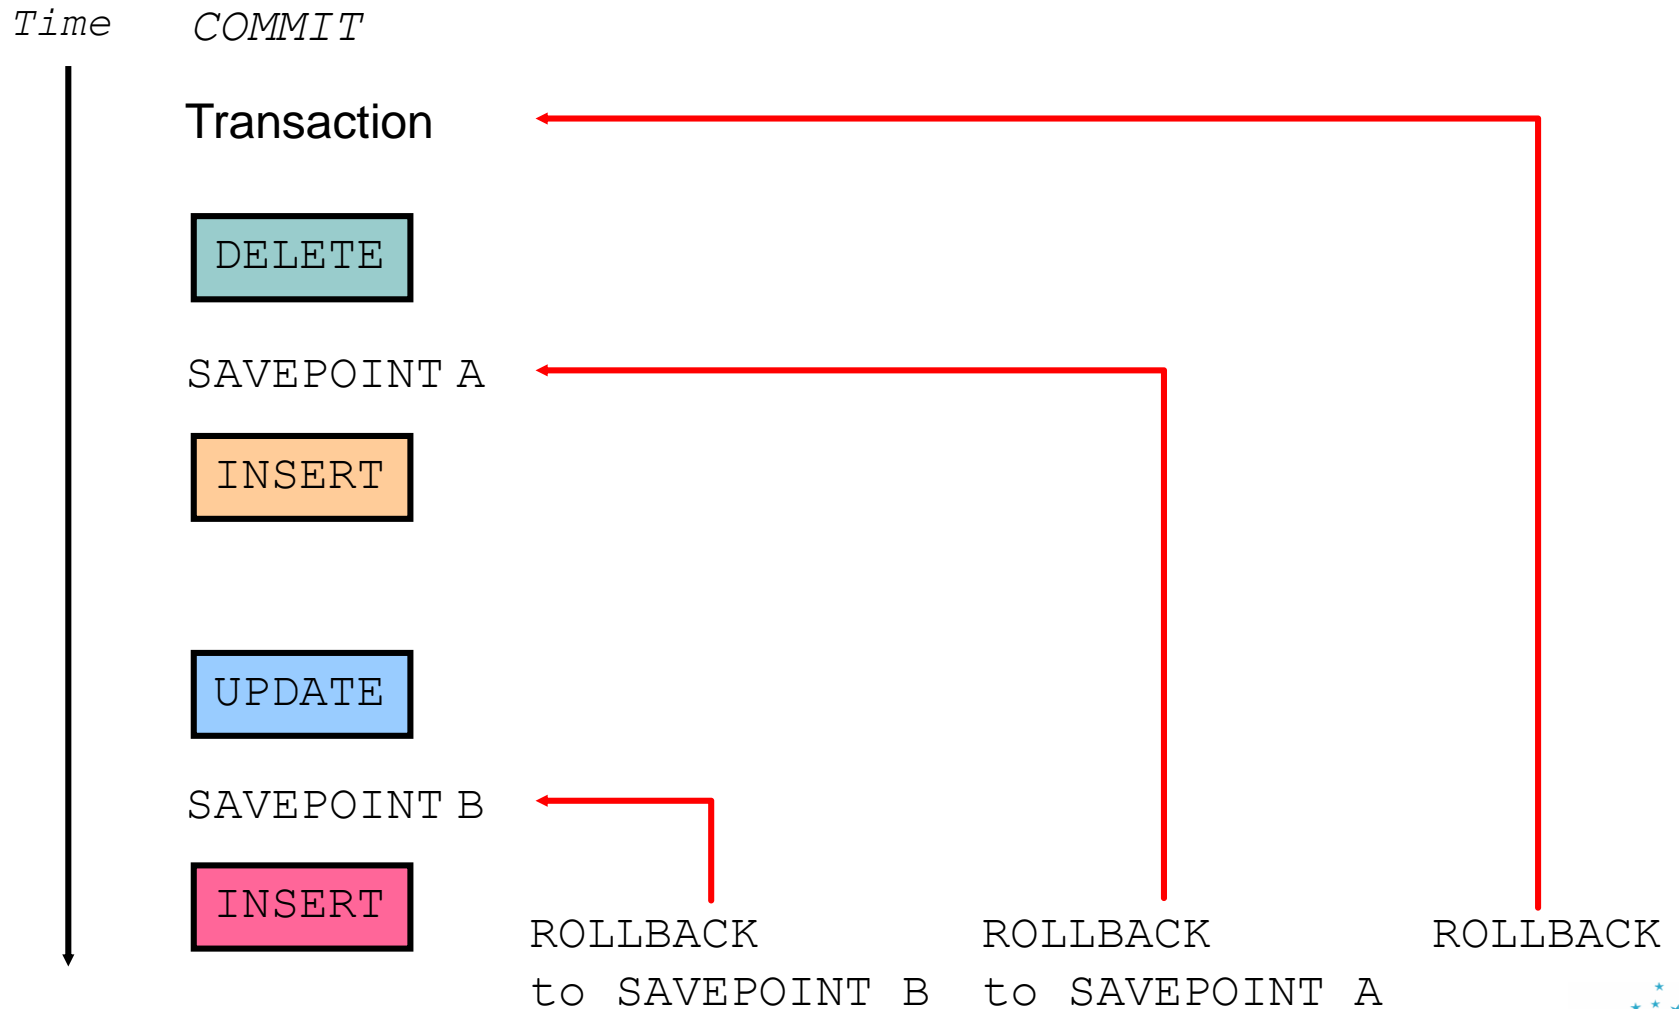
- Begin when the first DML SQL statement is executed
- End with one of the following events:
 - o A COMMIT or ROLLBACK statement is issued.
 - o A DDL or DCL statement executes (automatic commit).
 - o The user interact with SQL Developer or SQL*Plus.

Advantages of COMMIT and ROLLBACK Statements

With COMMIT and ROLLBACK statements, you can:

- Ensure data consistency
- Preview data changes before making changes permanent
- Group logically related operations

Controlling Transactions



Rolling Back Changes to a Marker

- Create a marker in a current transaction by using the `SAVEPOINT` statement.
- Roll back to that marker by using the `ROLLBACK TO SAVEPOINT` statement.

```
UPDATE...  
SAVEPOINT update_done;  
SAVEPOINT update_done succeeded  
INSERT...  
ROLLBACK TO update_done;  
ROLLBACK TO succeeded
```

Implicit Transaction Processing

- An automatic commit occurs under the following circumstances:
 - o DDL statement is issued
 - o DCL statement is issued
 - o Normal exit from SQL Developer or SQL*Plus, without explicitly issuing `COMMIT` or `ROLLBACK` statements
- An automatic rollback occurs under an abnormal termination of SQL Developer or SQL*Plus or a system failure.

State of the Data Before COMMIT or ROLLBACK

- The previous state of the data can be recovered.
- The current user can review the results of the DML operations by using the `SELECT` statement.
- Other users *cannot* view the results of the DML statements by the current user.
- The affected rows are *locked*; other users cannot change the data in the affected rows.

State of the Data After COMMIT

- Data changes are made permanent in the database.
- The previous state of the data is permanently lost.
- All users can view the results.
- Locks on the affected rows are released; those rows are available for other users to manipulate.
- All savepoints are erased.

Committing Data

- Make the changes:

```
DELETE FROM employees
WHERE  employee_id = 99999;
1 rows deleted

INSERT INTO departments
VALUES (290, 'Corporate Tax', NULL, 1700);
1 rows inserted
```

- Commit the changes:

```
COMMIT;
Commit complete
```


State of the Data After ROLLBACK

Discard all pending changes by using the ROLLBACK statement:

- Data changes are undone.
- Previous state of the data is restored.
- Locks on the affected rows are released.

```
DELETE FROM copy_emp;  
20 rows deleted  
ROLLBACK ;  
Rollback complete
```

State of the Data After ROLLBACK

```
DELETE FROM test;  
25,000 rows deleted
```

```
ROLLBACK;  
Rollback complete
```

```
DELETE FROM test WHERE id = 100;  
1 rows deleted
```

```
SELECT * FROM test WHERE id = 100;  
No rows selected
```

```
COMMIT;  
Commit complete
```

Statement-Level Rollback

- If a single DML statement fails during execution, only that statement is rolled back.
- The Oracle server implements an implicit savepoint.
- All other changes are retained.
- The user should terminate transactions explicitly by executing a `COMMIT` or `ROLLBACK` statement.

Read Consistency


- Read consistency guarantees a consistent view of the data at all times.
- Changes made by one user do not conflict with changes made by another user.
- Read consistency ensures that on the same data:
 - o Readers do not wait for writers
 - o Writers do not wait for readers

Implementation of Read Consistency


User A



```
UPDATE employees  
SET salary = 7000  
WHERE last_name = 'Grant';
```



```
SELECT *  
FROM userA.employees;
```



User B